

Lecture 1

Cezar Ionescu

Administrative matters

- course materials on GitHub
 - lectures notes will generally be available after the lecture, though drafts may appear before
 - exercises will also be posted on GitHub
- register for CATS points
 - assessment: homework to be handed in at the next meeting (or sent to me by email, or to PPWeekly)
 - every piece of homework is pass or fail
- course takes place Saturdays 10:00-12:30 at Ewert House
 - exceptions: **no class** on the **4th of May** and **15th of June!**
- main text: /Machine Learning/, Tom Mitchell, 1997
 - there are a couple of copies available in the ContEd library
 - you can buy it used for under 15 GBP on Amazon
 - but you should be able to complete the course using only the lecture notes

Introduction

- types of learning:
 - by rote
 - conditioning
 - from experience
 - any others?
- reasons for *machine learning*:
 - programming is hard, it would be better if computers learnt by themselves
 - to study human learning (and intelligence)
 - * perhaps we could then improve our own abilities to learn and to teach
- two main approaches:
 - modelling how we think and learn, without caring about the underlying physiological mechanisms
 - modelling the underlying physiological mechanism, without caring how they lead to thinking and learning

- **Definition** (Mitchell, p. 2) A computer program is said to **learn** from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .
 - how can we understand this mathematically?
 - * for example, by representing the various elements using sets and functions
 - * the most difficult to represent in this way appears to be *the task*
- **Note**
 - Sets and functions are the basic building blocks of mathematics. We assume them known.
 - The notation $f : \text{In} \rightarrow \text{Out}$ represents a function f taking inputs from the set In and “returning” values from the set Out . The notation $f(x)$ refers to **the** element of Out that f returns when given the input x (which, therefore, must be an element of In , i.e., $x \in \text{In}$).
 - In machine learning, we frequently encounter function-like “black boxes”, which, however, are not functions. The typical example is the Python “function” `Random.choice(seq)`, which takes as input a sequence (e.g., a list) and returns a randomly selected element of this sequence. Obviously, if the argument has more than one element, `Random.choice(seq)` will not be well-defined. We shall follow the generally accepted convention of using the function notation also for such black boxes. When we want to emphasise that we are not dealing with a proper function, we shall use a squiggly \rightarrow arrow instead of the standard straight one \rightarrow
- A preliminary mathematical representation:
 - the set of tasks: $\text{Task} = \text{In} \rightarrow \text{Out}$
 - the set of experiences: $\text{Experience} = \text{List } E$
 - measure of performance: $\text{perf} : (\text{Task}, \text{List } \text{Out}) \rightarrow \mathbb{R}$
 - * we expect some similarity between what we measure and the experience
 - machine learning system: $\text{learn} : \text{Experience} \rightarrow \text{Task}$
 - for all ex_1, ex_2, ins , we have

$\text{perf} (\text{learn } ex_1) ins \leq \text{perf} (\text{learn } (es_1 ++ es_2)) ins$

- Example: checkers learning
 - the task is to play a game: $\text{Play} = \text{Board} \rightarrow \text{Move}$
 - * a function in Play tells you how to move on the board (we assume that the board contains information about whose turn it is)
 - experience: $\text{Experience} = \text{Play} \rightarrow \text{List } \text{Game}$
 - * $\text{play} : (\text{Play}, \text{Play}) \rightarrow \text{Game}$
 - * the list of games is created by giving the play function the same argument *twice*
 - measure of performance: $\text{perf} : (\text{Play}, \text{List } \text{Play}) \rightarrow \mathbb{R}$

$\text{perf} (\text{learner}, [\text{adv}_1, \dots, \text{adv}_n]) =$
 $\text{percentage } [\text{score}(\text{play}(\text{learner}, \text{adv}_1)), \dots, \text{score}(\text{play}(\text{learner}, \text{adv}_n))]$

- Example: self-driving car
 - the task is to give steering commands based on sensor input: $\text{Drive} = \text{Sensor} \rightarrow \text{Steer}$

- the set of experiences: $\text{Experience} = \text{List}(\text{Sensor}, \text{Steer})$
- performance: $\text{perf} : (\text{Drive}, \text{Itinerary}) \rightarrow \text{Time}$
 - * $\text{perf}(\text{learner}, \text{itinerary})$ measures how long the learner drives along the given itinerary before making a mistake
- **Homework**
 - Give a similar interpretation for the handwriting recognition problem (Mitchell, page 3).

Concept learning

- idea: acquiring general concepts from examples
 - e.g., learn to recognise cats from images of animals
- what is a concept?
 - *nominalistic* view: the set of instances of the concept
- mathematically, we can identify a concept with a subset
 - e.g., X is the set of all images of animals, ' $C \subseteq X$ ' is the subset of images of cats
- subsets are in one-to-one correspondence with boolean-valued functions
 - $C \subseteq X$ can be replaced by $c : X \rightarrow \text{Bool}$ such that

$$\forall x \in X \quad c(x) = 1 \quad \text{iff} \quad x \in C$$

- Mitchell uses the functional view and defines:
 - **Concept learning:** inferring a boolean-valued function from training examples of input and output
- **Exercise:** Give an interpretation of concept learning as a learning task (i.e., identify the task, the experience, and the performance measure).
- **Notation:**
 - the training data $D = \{((x_1, c(x_1)) \dots, (x_n, c(x_n)))\}$
 - the subset of negative training examples $D_0 = \{(x, 0) \mid (x, 0) \in D\}$
 - the subset of positive training examples $D_1 = \{(x, 1) \mid (x, 1) \in D\}$
- **Example:** enjoyable day
 - we want to learn the concept $\text{enjoyable} : \text{Day} \rightarrow \{0, 1\}$
 - days are described via *attributes*:
 - * $\text{Day} = (\text{Sky}, \text{Temp}, \text{Humidity}, \text{Wind}, \text{Water}, \text{Forecast})$
 - $\text{Sky} = \{\text{Sunny}, \text{Cloudy}, \text{Rainy}\}$
 - $\text{Temp} = \{\text{Warm}, \text{Cold}\}$
 - $\text{Humidity} = \{\text{Normal}, \text{High}\}$
 - $\text{Wind} = \{\text{Strong}, \text{Weak}\}$
 - $\text{Water} = \{\text{Warm}, \text{Cool}\}$
 - $\text{Forecast} = \{\text{Same}, \text{Change}\}$
 - training data:

Nr	Sky	Temp	Humidity	Wind	Water	Forecast	Enjoyable
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

- **Two major related problems**

- **Day** contains 96 elements; there are 2^{96} concepts.
- We can represent a concept by a lookup table. In general, however, we are not going to be able to represent all concepts (the space could be infinite).
- The training data does not suffice to determine the concept we are looking for. It only fixes the values of the concept for 4 elements of its domain. Thus, there are 2^{92} remaining possibilities.
- The two problems force us to make two decisions:
 1. Choose a representation for **some** of the concepts. We are going to have to assume that the “real” concept can be represented that way. A representable concepts is called **hypothesis**.
 2. In general, we will still have many hypothesis consistent with the training data. The second decision is how to pick one of them.
- The assumptions under which we manage to learn the correct concept form **the inductive bias**.

Find-S

- Find-S solves problem 2 by choosing the *most specific* hypothesis that is consistent with the training data. Obviously, that implies that there exists an ordering from specific to general.
- Our hypothesis correspond to subsets. We have a natural ordering on subsets: \subseteq .

Find-S algorithm:

```
-- input: training data {(x1, c(x1)) ..., (xn, c(xn))}
--      hypothesis set H
h = min H -- set the current hypothesis h to "the" (or "a") smallest element of H
for i in 1:n
  if c(xi) = 0
    then keep h
  else if xi ∈ h then keep h
        else h = min {h' ∈ H | h ⊆ h' and xi ∈ h'}
-- output: "the" (or "a" most) specific hypothesis in H consistent with the training data
```

Remarks:

- If H contains all possible concepts, then the result of Find-S is D_1 .

- The result of Find-S can depend on arbitrary choices if the minimisation problems do not have a unique solution. Hence, the result might not be contained in C ! A simple example:

$$- X = \{a, b, c, d\}, H = \{\emptyset, \{a, b\}, \{a, c\}\}, D_1 = \{a\}$$

- The choice of H can avoid these problems.
- Hypothesis space for the weather example:
 - each hypothesis is described by a tuple of (Sky^* , $Temp^*$, $Humidity^*$, $Wind^*$, $Water^*$, $Forecast^*$) where $S^* = S \cup \{?, \emptyset\}$
 - notation: $s \sim s^*$ iff $s = s^*$ or $s^* = ?$ (s matches s^*)

Let h be described by $(s^*, t^*, u^*, wi^*, wa^*, f^*)$. Then

$$h(s, t, u, wi, wa, f) = s \sim s^* \text{ and } t \sim t^* \text{ and } u \sim u^* \text{ and } wi \sim wi^* \text{ and } wa \sim wa^* \text{ and } f \sim f^*$$