Lecture 1

Cezar Ionescu

11/05/2019

DEPARTMENT FOR CONTINUING EDUCATION

UNIVERSITY OF OXFORD

Administrative matters

- course materials on GitHub
  - lectures notes will generally be available after the lecture, though drafts may appear before
  - exercises will also be posted on GitHub

Administrative matters

- register for CATS points
  - assessment: homework to be handed in a the next meeting (or sent to me by email, or to PPWeekly)
  - every piece of homework is pass or fail

Administrative matters

- course takes place Saturdays 10:00-12:30 at Ewert House
  - exceptions: **no class** on the **4th of May** and **1st of June**!

Administrative matters

- main text: *Machine Learning*, Tom Mitchell, 1997
  - there are a couple of copies available in the ContEd library
  - you can buy it used for under 15 GBP on Amazon
  - but you should be able to complete the course using only the lecture notes

Types of learning

- types of learning:
  - by rote
  - conditioning
  - from experience
  - any others?

Why “machine learning”?

- reasons for *machine learning*:
  - programming is hard, it would be better if computers learnt by themselves
  - to study human learning (and intelligence)
    - perhaps we could then improve our own abilities to learn and to teach

Approaches to ML

- two main approaches:
  - modelling how we think and learn, without caring about the underlying physiological mechanisms
  - modelling the underlying physiological mechanism, without caring how they lead to thinking and learning

Definition

- **Definition** (Mitchell, p. 2) A computer program is said to **learn** from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E.

Sets and functions

Notation:

- `f : In → Out`
- `f(x)`
- `x ∈ In, f(x) ∈ Out`

Black boxes versus functions

- is `Random.choice(seq)` a function?
- we use ↠ instead of → when we need to distinguish black boxes from "real" functions

Mathematical representation

- the set of tasks: `Task`
- the set of experiences: `Experience`
- measure of performance: `perf : Task → ℝ`
- machine learning system: `learn : (Task, Experience) → Task`
  - learning means improving with experience (according to ‘`perf`):

```
for all t ∈ Task, for all e ∈ Experience:
    perf(t) ≤ perf(learn(t, e))
```

Example: checkers learning

- `Task = Board → Move`
- `experience : Task → List Game`
  - we assume `play : (Play, Play) ⇝ Game`
    - note the squiggly arrow in the type of `play`!

```
experience(t) = [play(t, t), play(t, t), ..., play(t, t)]
```

- `perf : (Task, List Task) → ℝ`
- we need a function `score : Game → {0, 1}`

```
perf (learner, [adv₁, ..., advₙ]) =
  (sum [score(play(learner, adv₁)), ...,
        score(play(learner, advₙ))]) * 100 / n
```

Example: self-driving car

- the task is to give steering commands based on sensor input:
  `Task = Sensor → Command`
- the set of experiences: `Experience = List (Sensor, Command)`
- performance: `perf : (Task, Itinerary) → Time`
  - `perf (learner, itinerary)` measures how long the learner drives along the given itinerary before making a mistake

Homework

- Give a similar interpretation for the handwriting recognition problem (Mitchell, page 3):
  - Task T: recognizing and classifying handwritten words within images
  - Performance measure P: percent of words correctly classified
  - Training experience E: a database of handwritten words with given classifications
- You have to fill in
  - `Task =`
  - `Experience =`
  - `perf :`
    - something about how `perf` is computed

Concept learning

- idea: acquiring general concepts from examples
  - e.g., learn to recognise cats from images of animals
- what is a concept?
  - *nominalistic* view: the set of instances of the concept

Mathematical representation of concepts

- mathematically, we can identify a concept with a subset
  - e.g., `X` is the set of all images of animals, '`C ⊆ X` is the subset of images of cats
- subsets are in one-to-one correspondence with boolean-valued functions
  - `C ⊆ X` can be replaced by `c : X → Bool` such that

```
∀ x ∈ X  c(x) = 1  iff x ∈ C
```

Definition

- Mitchell uses the functional view and defines:
  - **Concept learning:** inferring a boolean-valued function from training examples of input and output
- **Exercise:** Give an interpretation of concept learning as a learning task (i.e., identify the task, the experience, and the performance measure).

Notation

- the training data $D = \{(x_1, c(x_1)) \ldots, (x_n, c(x_n))\}$
- the subset of negative training examples
  $D_0 = \{(x, 0) \mid (x, 0) \in D\}$
- the subset of positive training examples
  $D_1 = \{(x, 1) \mid (x, 1) \in D\}$

Example

- we want to learn the concept `enjoyable : Day → {0, 1}`
- days are described via *attributes*:
  - `Day = (Sky, Temp, Humidity, Wind, Water, Forecast)`
    - `Sky = {Sunny, Cloudy, Rainy}`
    - `Temp = {Warm, Cold}`
    - `Humidity = {Normal, High}`
    - `Wind = {Strong, Weak}`
    - `Water = {Warm, Cool}`
    - `Forecast = {Same, Change}`

Training data

Nr	Sky	Temp	Humidity	Wind	Water	Forecast	Enjoyable
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

Two problems

- `Day` contains 96 elements; there are 2^{96} concepts.
- We can represent a concept by a lookup table, but not if it's too big!
- The training data does not suffice to determine the concept we are looking for.

Decisions

- The two problems force us to make two decisions:
  1. How to represent **some** of the concepts (the hypotheses space)
  2. Which hypothesis to pick.

Inductive bias

- The assumptions under which we manage to learn the correct concept form **the inductive bias**.

Hypothesis space for the weather example

- each hypothesis is described by a tuple of `(Sky*, Temp*, Humidity*, Wind*, Water*, Forecast*)`, where `S* = S ∪ {?, ∅}`
- notation: `s ~ s* iff s = s* or s* = ?` (s matches s*)

Let `h` be described by `(s*, t*, u*, wi*, wa*, f*)`. Then

```
h (s, t, u, wi, wa, f) = s ~ s* and t ~ t* and u ~ u* and
                         wi ~ wi* and wa ~ wa* and f ~ f*
```

Find-S

- `Find-S` solves problem 2 by choosing the *most specific* hypothesis that is consistent with the training data.
- Our hypothesis correspond to subsets. We have a natural ordering on subsets: $\subseteq$.

Find-S algorithm

```
-- input: training data {((x₁, c(x₁)) ..., (xₙ, c(xₙ))}
--        hypothesis set H
h = min H -- "the" (or "a") smallest element of H
for i in 1:n
  if c(xᵢ) = 0
    then keep h
    else if xᵢ ∈ h then keep h
                   else h = min {h' ∈ H | h ⊆ h' and xᵢ ∈ h'}
-- output: "the" (or "a") most specific hypothesis
--         consistent with D
```

Example

Work out how `Find-S` works on the weather example.

Remarks

- If H contains all possible concepts, then the result of Find-S is D_1.
- A bad situation for Find-S:
  - $X = \{a, b, c, d\}$, $H = \{\emptyset, \{a, b\}, \{a, c\}\}$, $D_1 = \{a\}$
- Another bad situation:

$X = \{a, b, c, d\}$, $H = \{\emptyset, \{a, b, c\}, \{a, b, d\}\}$,
$D_1 = \{a\}$, $D_0 = \{c\}$

- The choice of H can avoid these problems, as in the weather example.

Property of Find-S

If Find-S works, then

$$s = \text{Find-S}\ (D_0, D_1, H) \quad \text{implies}$$
$$D_0 \subseteq \neg s,\ D_1 \subseteq s, \text{ and}$$
$$\text{for all } h \in H,\ D_0 \subseteq \neg h \text{ and } D_1 \subseteq h \Rightarrow s \subseteq h$$

Exercise

What does `Find-S (D1, D0, H)` do?

A better Find-S

```
-- input: training data {((x₁, c(x₁)) ..., (xₙ, c(xₙ))}
--        hypothesis set H
S = allMin H -- start with all smallest element of H
repeat until S no longer changes:
  for i in 1:n
    if c(xᵢ) = 0
    then eliminate from S all {s | xᵢ ∈ s}
    else for all s ∈ S
        if xᵢ ∈ s
        then keep s
        else replace s with allMin {h' ∈ H | h ⊆ h' and xᵢ ∈ h'}
-- output: all most specific hyp consistent with D
```

Remarks

- why do we need to repeat the `for` loop?
- the algorithm terminates (why?)

Avoiding `repeat`

- we need to keep a record of the negative examples
- idea: do that in the same form as the record we keep for positive examples!
- this leads to the `Candidate-Elimination` algorithm